

Virtual Population analysis using eXtended Survivor Analysis

04 July, 2021

Required packages

To follow this tutorial you should have installed the following packages:

- FLR: FLCore, FLAssess, FLXSA, ggplotFL
- CRAN: [reshape] You can do so as follows,

```
install.packages(c("FLCore", "FLAssess", "FLXSA"), repos="http://flr-project.org/R")
```

```
# This chunk loads all necessary packages, trims pkg messages  
library(FLCore)  
library(FLAssess)  
library(FLXSA)
```

Introduction

What is VPA

Virtual population analysis (VPA) is a modeling technique commonly used in fisheries science for reconstructing the historical population structure of an age structured fish stock using information on the deaths of individuals in each time step. The time steps are typically, though not necessarily, annual and the deaths are usually partitioned into mortality due to fishing and natural mortality. In some instances natural mortality may be further partitioned into predation mortality and mortality from other causes, such as disease, senescence etc.

VPA is the most commonly used term to refer to cohort reconstruction techniques used in fisheries. It is virtual in the sense that the population size is not observed or measured directly but is inferred or backcalculated to have been a certain size in the past. Several different software implementations of cohort reconstruction for fish populations exist including ADAPT which is often used in Canada and the USA and XSA (Shepherd 1999) which is commonly used in Europe. The back-calculations in these implementations work the same way but they differ in the statistical methods used for “tuning” to indices of population size. Tuning refers to the use of auxiliary information to determine the terminal fishing mortalities and population numbers. Most tuning approaches involve a regression of fishing mortality against fishing effort to estimate population abundance at age through an iterative convergence to some threshold criterion. Relatively simple techniques, the Laurec-Shepherd method (???) for example, have been shown to work well with simulated data but there is little theoretical work to justify or validate these approaches (II and Deriso 1999).

A number of assessment methods are made available in FLR as well as the basic VPA tools to enable you to develop your own assessment methods. In this tutorial we will cover the basic VPA tools, simple methods for tuning a VPA and finally show how to run **FLXSA**.

Stock assessment methods within the FLR package structure

The package **FLAssess** contains the basic class for age and biomass based stock assessments. It provides a standard class, **FLAssess**, for data input, stock status estimation and diagnostic inspection. The **FLAssess** package has a variety of uses. It can be applied within a stock assessment working group setting or, alternatively, as part of the management procedure in a formal Management Strategy Evaluation (MSE). FLAssess provides a common interface for

existing stock assessment methods (e.g. XSA) allowing methods to be used interchangeably. It also includes various methods of general use such as setting up a short-term forecast (`stf`), running VPAs (`VPA` or `SepVPA`) and calculating F from catches. There are several steps to be completed when conducting an assessment. This tutorial considers only the process of running `VPA` and `FLXSA` stock assessment model. Additional tutorials are available that will introduce you to other parts of the **FLR** toolset.

We will start by importing the data sets for the North Sea Plaice stock and the fishery independent abundance indices. We will use these example data sets for all of the examples in this tutorial.

```
data(ple4)
data(ple4.indices)
```

The North Sea Plaice `FLStock` object already has values estimated for harvest and stock numbers. We should remove these first and replace them with `NA`.

```
harvest(ple4)[] <- NA
stock.n(ple4)[] <- NA
```

We should note at this point that the example below should not be considered the definitive assessment for the North Sea Plaice. We provide this example merely to show the procedure for conducting assessments using `FLR`.

The VPA method

The `VPA` method implements Pope's Virtual Population Analysis (`VPA`). It is called with the command `VPA` which returns an object of class `FLVPA` that is itself an extension of the `FLAssess` class. The `VPA` method estimates population numbers and fishing mortalities at age by back-calculating values down each cohort. To do this, the method requires initial values of harvest for the terminal age and terminal year in the `FLStock` object. These terminal values must be specified by the user prior to running the `VPA`. The arguments to the `VPA` method are the `FLStock` object for which values are to be calculated and two optional arguments.

The `range` method will show details of the age and year range of the `ple4` `FLStock` object. We can use this information to manually specify the terminal values in the harvest slot. In this instance we will set these values to 1.0. Remember to convert the values to be of type character when indexing the `FLQuants`.

```
harvest(ple4)[ac(range(ple4)["max"]), ] <- 1
harvest(ple4)[, ac(range(ple4)["maxyear"])] <- 1

ple4.vpa <- VPA(ple4, fratio = 1, fit.plusgroup = T)
ple4.new <- ple4 + ple4.vpa

## Have a look in stock number ##
stock.n(ple4.vpa)[, ac(2005:range(ple4)["maxyear"])]
```

An object of class "FLQuant"

, , unit = unique, season = all, area = unique

	year						
age	2005	2006	2007	2008	2009	2010	2011
1	727422.7	714089.7	1080988.1	860680.4	811461.7	1061031.1	1148353.6
2	740467.9	494608.8	505294.6	796821.9	631147.0	575619.3	762530.4
3	190271.1	413340.6	281588.2	289082.4	475010.1	402704.6	386692.2
4	273062.5	92337.1	200594.2	134282.1	147948.4	272930.1	245506.7
5	45186.6	135618.8	47632.4	106297.1	72270.2	80610.3	150395.0
6	34257.0	22604.0	72879.2	27001.3	61282.3	41221.8	44890.0
7	15049.4	17133.2	12161.7	40948.6	15756.0	36685.1	24284.4
8	7415.5	8112.3	9645.0	6851.3	23432.2	9331.3	22231.0

9	8478.0	4526.7	5157.5	6235.4	4486.5	15571.3	6298.0
10	5270.4	9513.7	10261.4	11855.6	14230.4	14703.3	23426.0

	year						
age	2012	2013	2014	2015	2016	2017	
1	889512.4	971160.1	1024564.4	488076.4	396779.2	189714.5	
2	883818.9	711636.0	755794.7	728978.7	320424.8	236873.2	
3	523353.9	604586.3	475967.2	484846.7	444943.1	182714.8	
4	230998.4	299022.9	341507.5	272373.6	268931.6	217290.6	
5	135024.8	124247.2	157044.0	176299.0	138428.9	124591.1	
6	80728.3	69083.3	61061.0	76403.8	84867.4	63401.4	
7	24682.3	41024.4	33728.5	28953.6	34952.6	35366.6	
8	14657.1	14415.4	22299.4	16289.6	11853.3	12982.4	
9	15148.4	9809.2	8957.9	12073.5	7364.8	4233.0	
10	22469.5	27525.0	26113.7	23046.2	21101.0	14646.2	

units: 1000

```
## Have a look in fishing mortality ##
harvest(ple4.vpa[, ac(2004:range(ple4)["maxyear"])]
```

An object of class "FLQuant"

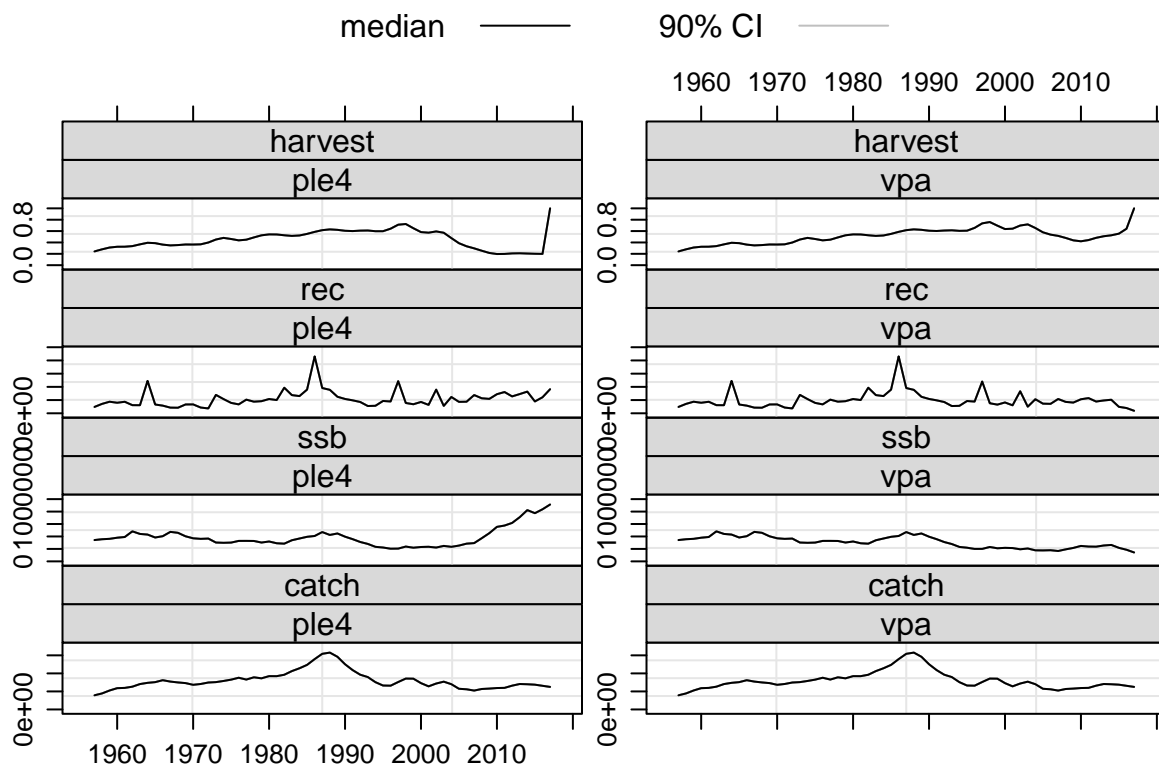
, , unit = unique, season = all, area = unique

	year								
age	2004	2005	2006	2007	2008	2009	2010	2011	2012
1	0.26372	0.28574	0.24587	0.20500	0.21018	0.24339	0.23035	0.16183	0.12311
2	0.54862	0.48301	0.46332	0.45843	0.41730	0.34934	0.29782	0.27638	0.27971
3	0.73103	0.62300	0.62299	0.64050	0.56985	0.45412	0.39488	0.41522	0.45974
4	0.64361	0.59985	0.56193	0.53505	0.51953	0.50724	0.49595	0.49787	0.52014
5	0.64425	0.59268	0.52104	0.46763	0.45075	0.46144	0.48541	0.52218	0.57014
6	0.70361	0.59288	0.51983	0.47648	0.43866	0.41312	0.42913	0.49813	0.57692
7	0.65271	0.51795	0.47458	0.47385	0.45821	0.42384	0.40088	0.40491	0.43779
8	0.48009	0.39359	0.35292	0.33619	0.32338	0.30868	0.29314	0.28359	0.30160
9	0.29044	0.26820	0.21355	0.16280	0.14003	0.14135	0.15646	0.17980	0.21238
10	0.29044	0.26820	0.21355	0.16280	0.14003	0.14135	0.15646	0.17980	0.21238


	year				
age	2013	2014	2015	2016	2017
1	0.15072	0.24038	0.32082	0.41586	1.00000
2	0.30222	0.34394	0.39370	0.46172	1.00000
3	0.47117	0.45817	0.48938	0.61671	1.00000
4	0.54399	0.56119	0.57682	0.66942	1.00000
5	0.61040	0.62049	0.63109	0.68087	1.00000
6	0.61697	0.64618	0.68204	0.77532	1.00000
7	0.50961	0.62781	0.79309	0.89040	1.00000
8	0.37576	0.51355	0.69381	0.92970	1.00000
9	0.25745	0.31989	0.40944	0.56452	1.00000
10	0.25745	0.31989	0.40944	0.56452	1.00000

units: f

```
## Plot results ##
plot(FLStocks(ple4=ple4, vpa=ple4.new))
```



```
plot(FLStocks(vpa=ple4.new))
```



The estimated population numbers and fishing mortality values at age from the VPA are now available in the returned object. Note that the terminal values for fishing mortality are the user defined values that were specified prior to running the VPA.

A simple method for tuning a VPA

As noted above the VPA method requires user defined terminal estimates of fishing mortality. This dependency limits the usefulness of the method since it is often the most recent, terminal, estimates that are of most concern to fishery managers. Additional catch at age and effort information, derived either from a sub component of the fishery or from a fishery independent source such as a research survey, can be used to ‘tune’ the assessment, as described above, and thereby obtain better estimates of fishing mortality and stock numbers in the most recent years. Several so-called ad hoc techniques for tuning a VPA have been developed. A relatively simple technique that has been widely used is the Laurec Shepherd method. This method can be easily implemented in FLR using the basic VPA tools that are provided in the **FLAssess** package.

The example shown below is a simple implementation that allows for a single tuning fleet. With a little extra effort it could be easily extended to accomodate multiple tuning fleets. The technical details of the method are not explained here.

```
# Define Laurec-Shepherd function #

lsm <- function(stock, index, fratio = 1, fit.plusgroup = T) {
  harvest(stock)[, ac(range(stock)["maxyear"])] <- 0.5
  diff <- 1
  while (diff > 1e-06) {
    stock <- stock + VPA(stock, fratio = fratio)
    ages <- range(index)["min"]:range(index)["max"]
    yrs <- range(index)["minyear"]:range(index)["maxyear"]
    stk <- trim(stock, year = yrs, age = ages)
    Cp <- catch.n(index)/catch.n(stk)
    q <- sweep(Cp * harvest(stk), 2, effort(index), "/")
    gmq <- apply(q, 1, function(x) exp(mean(log(x), na.rm = T)))
    mFp <- gmq * c(apply(effort(index), 1, mean))
    Fr <- mFp * (apply(Cp, 1, mean, na.rm = T))^-1
    Fnew <- c(Fr, rep(Fr[ac(max(ages)), ], 1))
    diff <- sum(abs(harvest(stock)[, ac(range(stock)["maxyear"])] -
                    Fnew))
    harvest(stock)[, ac(range(stock)["maxyear"])] <- c(Fnew)
  }
  res <- VPA(stock, fratio = fratio, fit.plusgroup = fit.plusgroup)
  index.res(res) <- FLQuants(q)
  return(res)
}
```

The new Laurec-Shepherd function can now be called without having to specify terminal values in the harvest slot. The arguments to the VPA method are also formally declared as arguments to our new function. Note that the function returns an object of class **FLVPA** that has been created from a call to the VPA method and that the catchability residuals are stored in the `index.res` slot of the returned object.

```
harvest(ple4)[] <- NA
stock.n(ple4)[] <- NA

ple4.LSvpa <- lsm(ple4, ple4.indices[[1]], fratio = 1, fit.plusgroup = T)

ple4.new2 <- ple4 + ple4.LSvpa

stock.n(ple4.LSvpa)[, ac(2005:range(ple4)["maxyear"])]
```

An object of class "FLQuant"
, , unit = unique, season = all, area = unique

	year						
age	2005	2006	2007	2008	2009	2010	2011
1	731905.9	719384.2	1090813.8	869716.2	820615.7	1089763.0	1181196.2
2	745531.6	498662.7	510082.8	805709.5	639320.0	583898.2	788517.0
3	191861.1	417913.5	285249.8	293407.5	483040.5	410092.4	394177.8
4	276099.6	93771.2	204718.7	137584.1	151851.6	280184.0	252183.0
5	45872.9	138358.7	48926.7	110020.5	75251.5	84134.6	156945.6
6	35005.9	23223.2	75353.0	28170.3	64645.8	43914.8	48073.1
7	15524.6	17808.8	12720.8	43182.9	16812.2	39724.5	26717.7
8	7719.9	8541.4	10255.2	7356.3	25450.6	10285.7	24977.9
9	8872.3	4801.7	5545.4	6787.2	4943.0	17396.3	7161.0
10	5515.5	10091.8	11033.1	12904.6	15678.6	16426.6	26636.0

	year					
age	2012	2013	2014	2015	2016	2017
1	935559.6	1043879.5	1156893.0	589017.9	636785.6	673808.7
2	913529.8	753296.2	821582.4	848664.4	411697.5	453860.7
3	546853.1	631453.0	513636.1	544320.8	553122.8	265193.9
4	237762.2	320250.5	365774.9	306403.6	322653.2	314942.5
5	141053.7	130354.0	176208.2	198204.4	169147.4	173058.4
6	86642.8	74524.5	66570.5	93696.5	104632.0	91114.7
7	27557.0	46362.5	38637.8	33923.4	50553.7	53183.4
8	16856.1	17012.9	27120.6	20719.6	16331.3	27043.7
9	17632.4	11797.7	11305.9	16428.4	11361.6	8266.6
10	26154.0	33104.8	32958.6	31359.0	32552.1	28602.4

units: 1000

```
harvest(ple4.LSvpa[, ac(2004:range(ple4)["maxyear"])]]
```

An object of class "FLQuant"

```
, , unit = unique, season = all, area = unique
```

	year								
age	2004	2005	2006	2007	2008	2009	2010	2011	2012
1	0.26215	0.28372	0.24382	0.20296	0.20776	0.24033	0.22356	0.15697	0.11669
2	0.54513	0.47882	0.45856	0.45301	0.41162	0.34402	0.29292	0.26597	0.26929
3	0.72533	0.61591	0.61364	0.62913	0.55866	0.44465	0.38623	0.40553	0.43508
4	0.63651	0.59091	0.55053	0.52097	0.50340	0.49049	0.47955	0.48101	0.50102
5	0.63407	0.58072	0.50767	0.45205	0.43174	0.43858	0.45970	0.49411	0.53801
6	0.68809	0.57582	0.50191	0.45674	0.41616	0.38695	0.39693	0.45647	0.52530
7	0.63370	0.49750	0.45191	0.44768	0.42871	0.39135	0.36398	0.36061	0.38229
8	0.46306	0.37482	0.33195	0.31276	0.29758	0.28048	0.26210	0.24825	0.25681
9	0.27920	0.25466	0.20003	0.15053	0.12790	0.12744	0.13887	0.15637	0.17965
10	0.27920	0.25466	0.20003	0.15053	0.12790	0.12744	0.13887	0.15637	0.17965

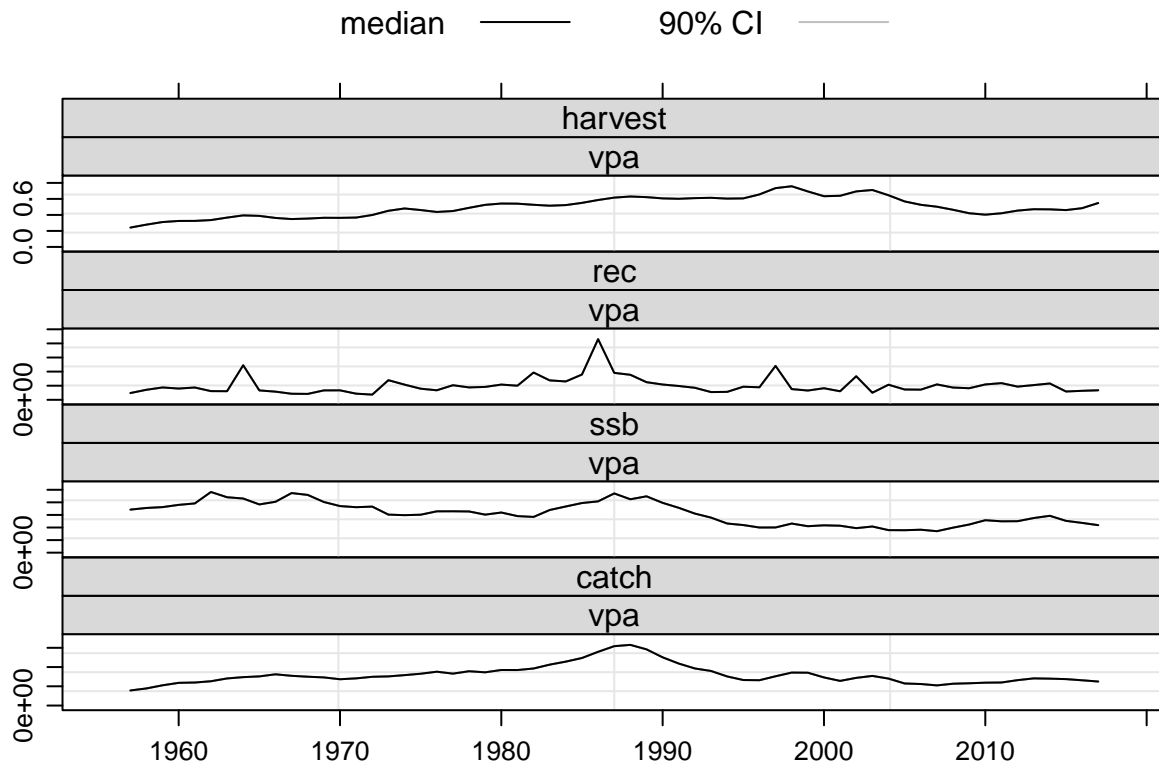
	year				
age	2013	2014	2015	2016	2017
1	0.13947	0.20983	0.25817	0.23864	0.19741
2	0.28294	0.31169	0.32808	0.33983	0.40273
3	0.44601	0.41661	0.42296	0.46319	0.57450
4	0.49744	0.51272	0.49413	0.52295	0.57557
5	0.57199	0.53160	0.53885	0.51865	0.60975
6	0.55690	0.57416	0.51703	0.57670	0.58221
7	0.43620	0.52315	0.63102	0.52558	0.54799
8	0.30864	0.40128	0.50084	0.58086	0.36387

```
9 0.20924 0.24468 0.28392 0.32874 0.39346
10 0.20924 0.24468 0.28392 0.32874 0.39346
```

```
units: f
```

```
# Compare the results with previous fits.
```

```
plot(FLStocks(vpa=ple4.new2))
```



FLXSA

The Laurec-Shepherd method above is a relatively simple technique for tuning a VPA. XSA is a more sophisticated method that uses information on individual cohort sizes to estimate survivors at each age in the terminal population. Although the modelling approach is more involved the method requires the same input of catch numbers at age and indices of catch per unit effort and it retains at its core the basic VPA method. The details of the XSA method are too complex to show here, or to code individually as we have for the Laurec-Shepherd approach. Instead the FLXSA method has been developed as an additional package to FLAssess.

The FLXSA control object

The `FLXSA.control` object contains all of the user defined model settings for running an XSA analysis. It can be created in several different ways. The simplest method is to accept all of the default settings by calling the `FLXSA.control` function without any extra arguments:

```
FLXSA.control()
```

```
tol          1e-09
maxit        30
min.nse      0.3
fse          0.5
rage         0
```

```

qage          10
shk.n         TRUE
shk.f         TRUE
shk.yrs       5
shk.ages      5
window        100
tsrange       20
tspower       3
vpa           FALSE

```

Alternatively the default settings can be over-written by specifying values at the point of creation or by overwriting them afterwards.

```

ctrl <- FLXSA.control(maxit = 50, qage = 8)
ctrl <- FLXSA.control()
slot(ctrl, 'qage') <- as.integer(8)
slot(ctrl, 'maxit') <- as.integer(50)

```

Note that in the example above, when modifying the control object after creation, it is necessary to coerce the values 8 and 50 to type integer. This is because the default type numeric cannot be used in this slot. Such coercion is not necessary when using the `FLXSA.control` function as this check is performed internally by the function. You can use the `getSlots` function to determine the class of object associated with any given slot.

```

xsa.control <- FLXSA.control(maxit = 50, fse = 2.5)
ple4.xsa <- FLXSA(ple4, ple4.indices, xsa.control)
ple4.xsa.t1 <- FLXSA(ple4, ple4.indices[[1]], xsa.control)

```

Once the control object has been created, the XSA analysis can be run as a one-line command. The “`FLXSA`” function returns an object of class `FLXSA` which extends the `FLAssess` class. The `FLXSA` object contains all of the information in the `FLAssess` class plus additional information specific to the XSA assessment method, such as the survivors estimates and their internal and external standard errors. The control object used for the assessment is also stored in the returned `FLXSA` object to provide a record of what settings were used for that particular run. All of the settings in the returned control object will remain the same except for the `maxit` slot that contains the maximum number of iterations for the analysis. This value will be overwritten with the actual number of iterations taken to reach convergence, if indeed the model had converged before the maximum number initially specified.

XSA Results

Appart from the model diagnostics, the `FLXSA` method returns two important results, namely the estimated values of fishing mortality and population numbers at age. These are returned as `FLQuants` and are stored in the `harvest` and `stock.n` slots, respectively, of the `FLXSA` object. These estimated values can be very easily read back into an `FLStock` object using the `+` operator. Once the results have been read back into a `FLStock` object we can look at some of the key information such as SSB, recruitment and mean fishing mortality values. But before concentrating too much on the results of the assessment it is advisable to first investigate some of the model diagnostics.

```

ple4.new <- ple4 + ple4.xsa
ple4.ssb <- ssb(ple4.new)
ple4.rec <- rec(ple4.new)
ple4.fbar <- fbar(ple4.new)

```

XSA Diagnostics

There are many diagnostic checks that one might be interested in conducting to examine the model fit. The first might be to see if the model has reached convergence within the specified number of iterations.


```
slot(slot(ple4.xsa, "control"), "maxit")
```

```
[1] 50
```

Additionally one can check for discrepancies between the internal and external standard errors of the survivors estimates. Very often plots of the catchability residuals are made to inspect for any obvious trends or departures from the assumption of constant catchability over time. Some examples of these plots and details of their creation from FLR objects are provided below but you should also consult the tutorial on lattice plotting and advanced graphics for FLR to see examples of other ways to graphically display your data.

There are several ways to access diagnostic information about your fitted XSA model. The easiest is perhaps to use the `diagnostics` function, which will replicate the diagnostic output produced by the original VPA suite (developed in the early 1990's). Note that this function merely outputs the results to the screen and no object is created by the method. The function was created to allow the user to cut and paste the information from the console to a report. The output can be quite long, particularly if the assessment comprises a large number of ages and many tuning indices. The standard output can be divided roughly into eight sections each providing different information about the model and the fit. These sections comprise the model dimensions; parameter settings; regression weights; the estimated fishing mortalities and population numbers for the last 10 years; the aggregated survivors estimates; the log catchability residuals for each of the tuning indices and finally the individual survivors estimates for each year-class represented in the terminal year.

In order to make this document more readable we will print out only a few sections of the diagnostic output at a time. We can do this by passing a vector of TRUE and FALSE values to the sections argument of the `diagnostics` method. By default all sections are set to TRUE so that all of the information is output to the screen. In order to reduce the quantity of output further

we will run a new XSA for a reduced number of ages and with only one tuning index and will start by outputting only the dimension information and the parameter settings from our diagnostics.

```
ple4.xsa2 <- FLXSA(trim(ple4, age = 1:7), ple4.indices[[3]],
  xsa.control)
diagnostics(ple4.xsa2, sections = c(T, T, rep(F, 6)))
```

```
FLR XSA Diagnostics 2021-07-04 22:04:53
```

```
CPUE data from indices
```

```
Catch data for 61 years 1957 to 2017. Ages 1 to 7.
```

```
fleet first age last age first year last year alpha beta
1  SNS          1          6      1970      2017 <NA> <NA>
```

```
Time series weights :
```

```
  Tapered time weighting applied
  Power =   3 over 20 years
```

```
Catchability analysis :
```

```
  Catchability independent of size for all ages
```

```
  Catchability independent of age for ages > 6
```

```
Terminal population estimation :
```

```
  Survivor estimates shrunk towards the mean F
```

of the final 5 years or the 5 oldest ages.

S.E. of the mean to which the estimates are shrunk = 2.5

Minimum standard error for population
estimates derived from each fleet = 0.3

prior weighting not applied

Next we can output the regression weights and the fishing mortalities and population numbers for the last 10 years and also the aggregated survivors estimates.

```
diagnostics(ple4.xsa2, sections = c(F, F, T, T, T, T, F, F))
```

Regression weights

	year										
age	2008	2009	2010	2011	2012	2013	2014	2015	2016	2017	
all	0.751	0.82	0.877	0.921	0.954	0.976	0.99	0.997	1	1	

Fishing mortalities

	year										
age	2008	2009	2010	2011	2012	2013	2014	2015	2016	2017	
1	0.207	0.247	0.200	0.124	0.105	0.104	0.147	0.122	0.183	0.107	
2	0.443	0.343	0.304	0.231	0.204	0.249	0.218	0.211	0.136	0.284	
3	0.535	0.499	0.385	0.429	0.359	0.306	0.349	0.261	0.254	0.171	
4	0.550	0.456	0.580	0.479	0.549	0.372	0.295	0.380	0.262	0.242	
5	0.570	0.506	0.410	0.689	0.535	0.674	0.340	0.240	0.346	0.223	
6	0.692	0.609	0.498	0.383	1.007	0.552	0.786	0.265	0.182	0.312	
7	0.692	0.609	0.498	0.383	1.007	0.552	0.786	0.265	0.182	0.312	

XSA population number (Thousand)

	age								
year		1	2	3	4	5	6	7	
2008		873829	762854	304706	129307	89155	19273	30076	
2009		801925	642724	443318	161500	67535	45627	11911	
2010		1207505	566603	412620	243597	92626	36835	30880	
2011		1468479	894604	378145	254039	123415	55638	25433	
2012		1036439	1173206	642406	222804	142348	56052	13725	
2013		1382295	844446	865899	406006	116431	75436	38620	
2014		1603721	1127599	595646	577062	253287	53708	28858	
2015		1171171	1252542	820617	379995	388764	163122	68342	
2016		810347	938140	917890	571988	235246	276692	124405	
2017		1195955	610611	741191	644206	398100	150631	83772	

Estimated population abundance at 1st Jan 2018

	age							
year		1	2	3	4	5	6	7
2018		0	973125	416096	565666	457969	288826	99986

And finally we can output the catchability residuals and the individual survivors estimates. Note that very little thought went into the parameter settings for this particular model fit so please don't interrogate the output presented here too closely. Also note that we do not normally expect the diagnostics output to be broken up as we have here. We present it in this way purely to make it more presentable in this document.

By default all sections are set to TRUE so it is very likely that you won't need to give this argument at all when calling the diagnostics method.

```
diagnostics(ple4.xsa2, sections = c(F, F, F, F, F, F, T, T))
```

Fleet: SNS

Log catchability residuals.

year												
age	1970	1971	1972	1973	1974	1975	1976	1977	1978	1979	1980	
1	0.113	0.922	1.057	1.332	0.162	0.239	0.132	0.504	0.375	0.715	1.506	
2	1.929	2.996	2.471	1.908	1.773	1.370	0.527	1.119	2.057	1.965	2.128	
3	2.505	1.262	2.440	1.875	1.747	2.565	0.449	1.537	0.203	1.412	0.461	
4	2.032	0.150	-0.335	1.402	0.867	0.156	0.084	1.271	0.224	0.742	0.603	
5	1.015	0.223	0.938	-1.072	0.773	-0.406	NA	1.019	-0.233	-0.730	-0.013	
6	-0.364	-0.634	NA	-1.346	-1.382	NA	-1.832	NA	-1.264	0.206	-0.981	
year												
age	1981	1982	1983	1984	1985	1986	1987	1988	1989	1990	1991	
1	0.677	1.062	0.709	1.068	0.283	0.598	0.360	0.503	0.774	0.601	1.535	
2	2.248	1.485	1.299	1.390	1.773	0.871	1.462	1.418	1.655	1.101	1.644	
3	0.896	1.009	-0.655	0.946	0.895	0.689	0.488	2.003	1.669	1.332	0.891	
4	-0.702	0.035	-0.523	-1.132	0.217	-0.210	0.250	1.910	2.710	1.285	1.529	
5	0.032	-1.108	-1.516	NA	-0.925	-1.052	-0.653	0.559	1.490	0.998	0.538	
6	-1.114	-0.759	-2.432	NA	-2.783	-1.378	-1.131	-1.053	1.306	0.184	-0.270	
year												
age	1992	1993	1994	1995	1996	1997	1998	1999	2000	2001	2002	
1	1.481	1.164	0.973	0.739	0.469	1.101	1.119	1.289	0.770	0.433	0.425	
2	1.977	1.774	1.513	0.859	1.694	0.613	1.815	1.706	0.462	0.513	0.003	
3	1.449	0.786	0.780	0.371	1.489	0.319	2.106	2.528	0.567	-0.096	-0.351	
4	1.257	-0.400	-0.478	-0.056	-0.389	-0.676	-0.350	1.683	1.319	0.171	-0.554	
5	0.811	0.024	-0.746	-0.041	0.386	NA	-0.613	0.081	-1.638	2.240	0.227	
6	0.111	0.191	-2.608	NA	NA	-0.322	NA	NA	-1.035	0.315	-0.209	
year												
age	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013	
1	NA	0.391	0.183	0.316	0.100	0.310	0.435	-0.226	-0.046	-0.847	-0.176	
2	NA	0.478	0.017	0.232	0.452	0.312	0.122	0.113	-0.856	-0.206	0.186	
3	NA	0.721	-0.240	-0.024	-0.691	0.323	-0.019	0.045	-0.378	-0.162	0.378	
4	NA	0.125	1.165	0.161	-0.650	0.883	-0.955	-0.909	-0.711	-0.442	0.486	
5	NA	-0.088	-0.421	0.641	NA	0.085	-0.213	-1.290	0.412	-0.533	0.853	
6	NA	0.052	0.028	NA	-0.200	-0.023	NA	NA	0.129	-0.063	0.049	
year												
age	2014	2015	2016	2017								
1	-0.176	0.086	0.021	0.006								
2	-0.417	0.276	-0.215	-0.031								
3	-0.376	0.539	-0.128	0.137								
4	0.148	0.668	0.264	0.154								
5	-0.587	0.152	0.837	-0.213								
6	-0.086	0.219	-0.068	0.009								

Mean log catchability and standard error of ages with catchability independent of year class strength and constant w.r.t. time

	1	2	3	4	5	6
Mean_Logq	-4.1782	-5.5729	-6.3623	-7.0081	-7.4844	-7.3835
S.E_Logq	0.5162	0.8819	0.8866	0.8817	0.8357	0.9003

Terminal year survivor and F summaries:

,Age 1 Year class =2016

```
source
      scaledWts survivors yrcls
SNS      0.977    978305 2016
fshk     0.023    773780 2016
```

,Age 2 Year class =2015

```
source
      scaledWts survivors yrcls
SNS      0.969    403220 2015
fshk     0.031    604524 2015
```

,Age 3 Year class =2014

```
source
      scaledWts survivors yrcls
SNS      0.971    648476 2014
fshk     0.029    293770 2014
```

,Age 4 Year class =2013

```
source
      scaledWts survivors yrcls
SNS      0.913    533680 2013
fshk     0.087    277605 2013
```

,Age 5 Year class =2012

```
source
      scaledWts survivors yrcls
SNS      0.906    233051 2012
fshk     0.094    134437 2012
```

,Age 6 Year class =2011

```
source
      scaledWts survivors yrcls
SNS      0.981    100592 2011
fshk     0.019    160164 2011
```

Remember that the diagnostics method will only output text to the console, enabling you to copy and paste the output to a report or other document. If you want to access the diagnostic data you will need to access the specific slots of the returned **FLXSA** object. The information that you will require is contained in various slots. The individual estimates of population number from each source (ie. tuning series and F shrinkage)

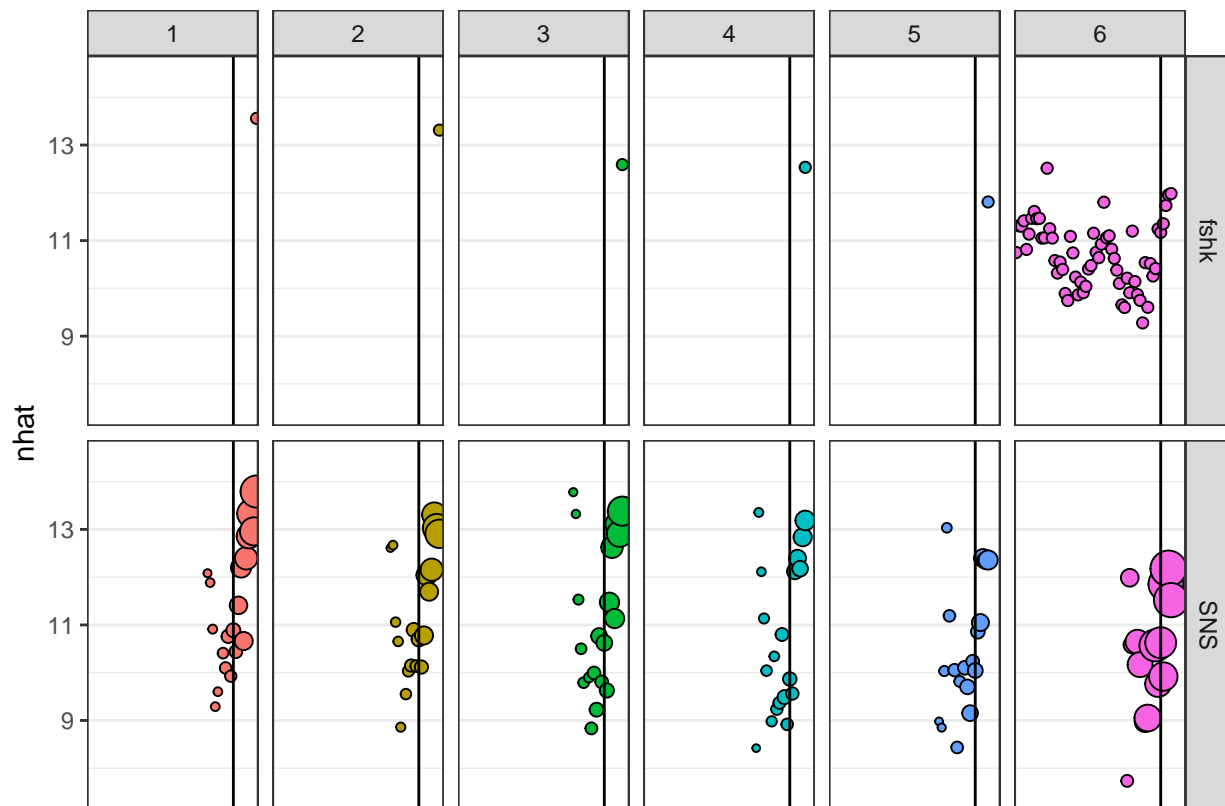
and their individual weightings are stored as a dataframe in the `diagnostics` slot of the returned object. Other slots contain the internal and external standard errors; the log catchability residuals. For a more thorough description of the XSA diagnostics you should consult the VPA users manual.

Plotting Diagnostics

Very often the quickest and simplest way to determine the fit of the model is through visual inspection of the various diagnostic outputs.

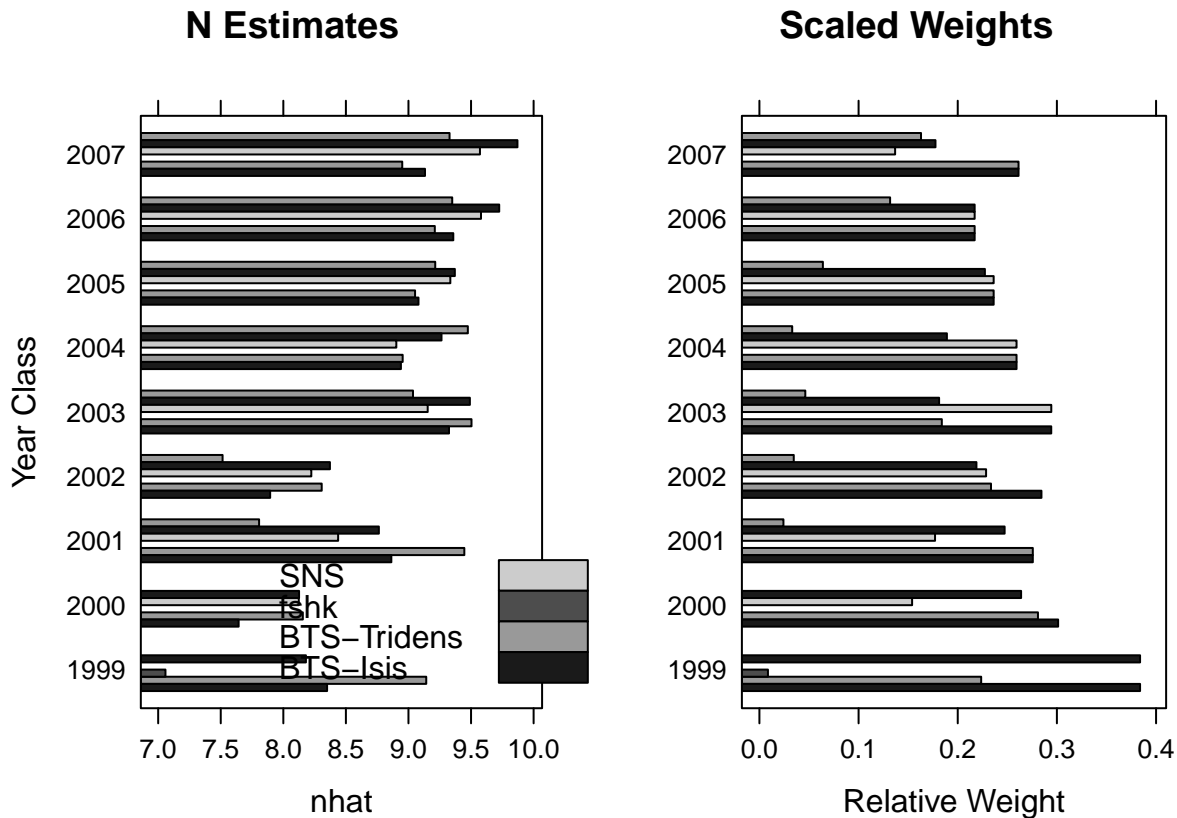
The default plot for "FLXSA" class shows the weight given to each of the indices, including the shrinkage, to estimate total numbers at age along ages and years. The size of the bubbles in the plot is proportional to the weight given to the index to estimate the terminal numbers at age. The rows corresponds with the indices used and the columns with age classes. The y axis represent the estimate of numbers at age obtained from each index.

```
plot(ple4.xsa2)
```



Below we provide examples of how to extract the relevant information from the return FLXSA object and to plot it using a variety of lattice functions available to R. We start by plotting the log catchability residuals at age from each of the three tuning series. The data are stored as an FLQuants object in the `index.res` slot of the FLXSA object. First we need to assign names to each of the FLQuant objects so we know which fleet they represent.

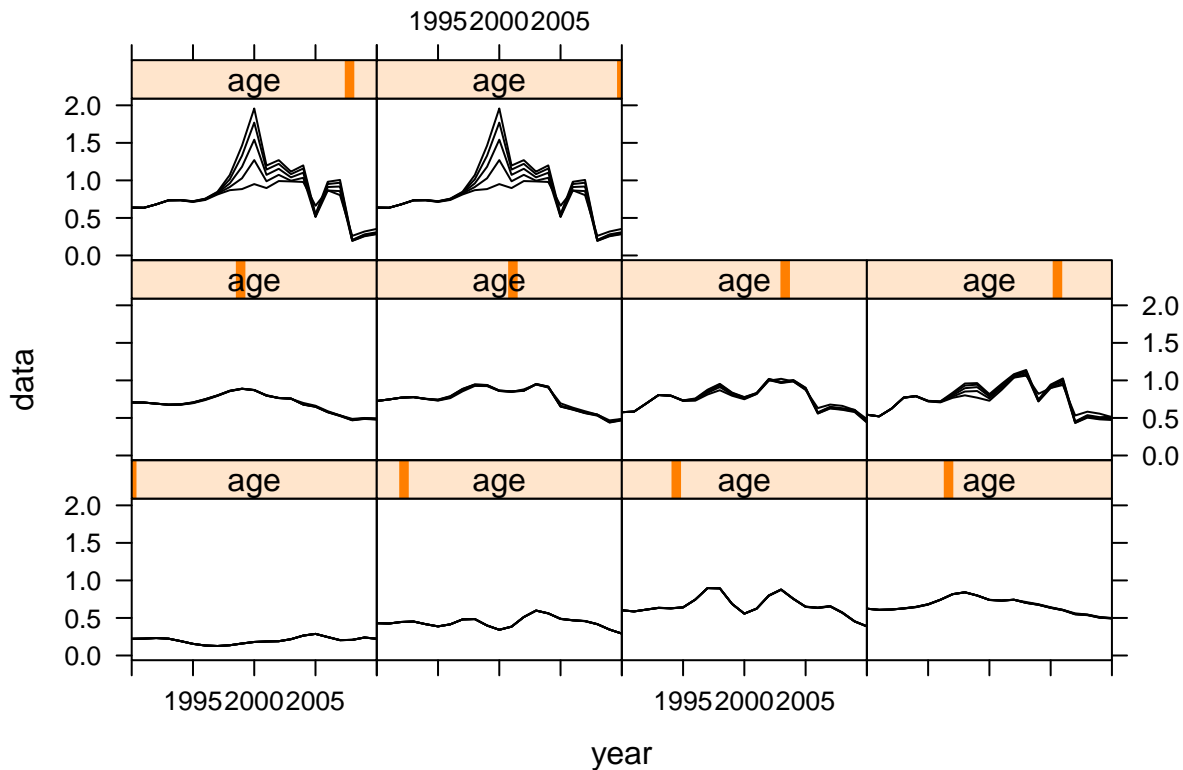
```
names(ple4.xsa@index.res) <- names(ple4.indices)
plot(xyplot(data ~ year | ac(age) + qname, data = index.res(ple4.xsa),
  panel = function(x, y, ...) {
    panel.xyplot(x, y, ...)
    panel.loess(x, y, ...)
    panel.abline(h = 0, col = "green", lty = 2)
  })))
```

Sensitivity to different model settings.

The simplified calling format of FLXSA makes it very easy to run multiple analyses to investigate model sensitivity to parameter settings. A wide variety of such investigations are possible. In this simple example we will look at the effect that different F shrinkage standard errors have on the terminal estimates of fishing mortality. We start by creating a vector of F shrinkage values to be used in the analyses and by creating an FLQuant with sufficient dimensions to store the results. To do this we use the propagate function to extend an FLQuant in the 6th dimension by the number of runs that we are going to perform. The estimates of fishing mortality for each XSA run are stored in the FLQuant using the 6th dimension to hold each iteration. The results show little sensitivity to increasing F shrinkage values at values between 1.0 and 2.5 .

```
fsevals <- seq(0.5, 2.5, by = 0.5)
res <- propagate(harvest(ple4), length(fsevals))
for (i in 1:length(fsevals)) {
  xsa.control <- FLXSA.control(fse = fsevals[i])
  iter(res, i) <- harvest(FLXSA(ple4, ple4.indices, xsa.control))
}
plot(xyplot(data ~ year | age, groups = iter, data = res, type = "l",
  col = "black", xlim = c(1990:2010)))
```

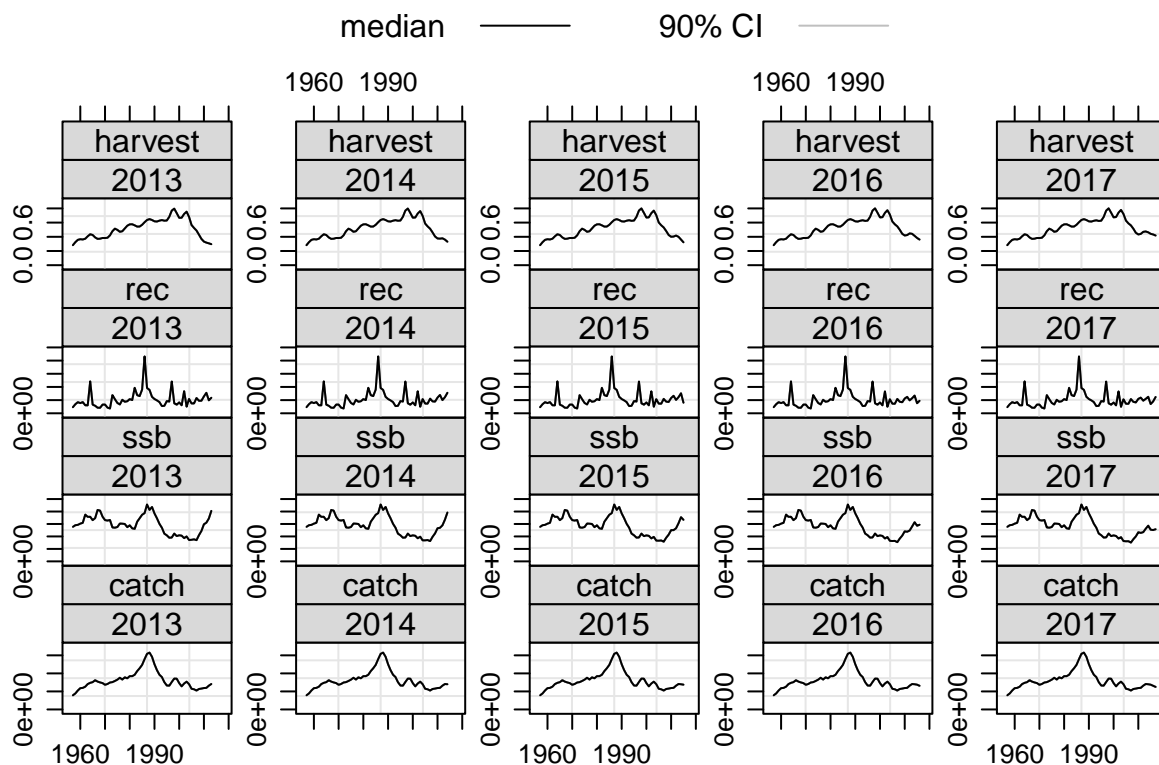


Retrospective Analyses

An important diagnostic check is to see how the estimated values vary as the time series of the input data changes. We can make use of existing R functions to apply the same assessment model to successively truncated the time series of input data. In this example we are using `window` to truncate the `FLStock` object to the specified year range, the `+` operator to pass the results of the XSA into the `FLStock` object and the `tapply` function to perform this action over the year range 2004:2008. Note that the resulting object, called `ple4.ret`, is of class `FLStocks` ie. a list of `FLStock` objects, each one having a separate year range.

```
retro.years <- 2013:2017
ple4.retro <- tapply(retro.years, 1:length(retro.years), function(x){
  window(ple4,end=x)+FLXSA(window(ple4,end=x),ple4.indices)
})

# coerce into FLStocks object
ple4.retro <- FLStocks(ple4.retro)
# full retrospective summary plot
ple4.retro@names=ac(c(retro.years))###Add years to legend
plot(ple4.retro)
```

More information

- You can submit bug reports, questions or suggestions on this tutorial at <https://github.com/flr/doc/issues>.
- Or send a pull request to <https://github.com/flr/doc/>
- For more information on the FLR Project for Quantitative Fisheries Science in R, visit the FLR webpage, <http://flr-project.org>.

Software Versions

- R version 4.1.0 (2021-05-18)
- FLCore: 2.6.16
- FLXSA: 2.6.4
- FLAssess: 2.6.3
- **Compiled:** Sun Jul 4 22:05:17 2021

License

This document is licensed under the Creative Commons Attribution-ShareAlike 4.0 International license.

Author information

Dorleta Garcia AZTI. Marine Reserach Unit. Txatxarramendi Ugartea z/g, 48395, Sukarrieta, Basque Country, Spain.

Alessandro MANNINI. European Commission, DG Joint Research Centre, Directorate D - Sustainable Resources, Unit D.02 Water and Marine Resources, Via E. Fermi 2749, 21027 Ispra VA, Italy. <https://ec.europa.eu/jrc/>

References

II, Terrance J. Quinn, and Richard B. Deriso. 1999. *Quantitative Fish Dynamics*. New York: Oxford University Press.

Shepherd, J. G. 1999. “Extended Survivors Analysis: An Improved Method for the Analysis of Catch-at-Age Data and Abundance Indices.” *ICES Journal of Marine Science* 56: 584–91.